Robust Medical Image Classification with Curvature Regularization on the PATHMNIST

Guohao Yang*, Jiacheng Qi and Xubin Sun
Department of Computer Science, Wenzhou Kean University, Wenzhou, Zhejiang, 325060, China
*Corresponding author's e-mail: yanggu@kean.edu

ABSTRACT

In medical image analysis, the robustness and accuracy of classification models are paramount, especially under adverse conditions such as data scarcity, class imbalance, and potential adversarial attacks. This paper introduces a novel deep learning architecture employing curvature regularization (CURE) to enhance the robustness of medical image classification on the PATHMNIST dataset. Our approach integrates advanced convolutional techniques including depthwise and dilated convolutions with an attention mechanism, focusing on precise and detailed feature extraction essential for medical diagnostics. We incorporate curvature regularization to stabilize the learning process by controlling the magnitude of the Hessian's eigenvalues, making the model less sensitive to input variations. The effectiveness of our architecture is demonstrated through rigorous testing, including adversarial scenarios using the Fast Gradient Sign Method (FGSM) and Projected Gradient Descent (PGD). Results show our model not only achieves superior accuracy compared to traditional architectures like ResNet-18 and ResNet-50 but also maintains higher resilience against adversarial attacks. This robustness is critical for practical deployment in medical image analysis

Keywords: Medical Image Classification, Curvature Regularization, Adversarial Robustness

1. INTRODUCTION

Deep learning has substantially improved medical image analysis, significantly enhancing classification and segmentation tasks across various domains such as breast, lesion, chest, and lung imaging. The success of these models often relies on substantial and well-labeled datasets. However, medical imaging frequently faces challenges like data scarcity and class imbalance, which can impede model effectiveness. Transfer learning has become a crucial technique, leveraging knowledge from related tasks to improve efficiency and performance in data-limited situations. Traditional transfer learning methods often employ pre-trained models from vast general image datasets, such as ImageNet, and fine-tune them for specific medical tasks. Despite their effectiveness, these models may struggle to adapt to the unique characteristics of medical images. Recent advancements have introduced models like OpenAI's CLIP [1], which, although trained on diverse internet data, do not inherently possess domain-specific knowledge, particularly in specialized fields such as medicine.

This paper introduces a robust algorithmic framework for the classification of medical images using the PATHMNIST dataset, emphasizing resilience against machine malfunctions or sensor anomalies. Our proposed model integrates advanced convolutional techniques with Curvature Regularization (CURE) to enhance robustness and adaptability. The architecture employs specialized convolutions including depthwise and dilated convolutions, tailored for detailed and precise feature extraction from medical images. Furthermore, the model includes attention mechanisms to improve interpretability and focus, which are essential for accurate medical diagnostics. The main contributions of this paper are as follows:

- We propose a robust deep learning architecture that combines advanced convolutional strategies and Curvature Regularization to enhance stability and performance under adverse conditions.
- Through experiments on the PATHMNIST dataset, our model not only achieves good accuracy but also demonstrates improved resilience and efficiency during robust test.

2. RELATED WORKS

In the realm of medical image analysis using machine learning, recent advancements have significantly leveraged the PathMNIST dataset for benchmarking and methodological innovation. [2, 3] have explored the realms of privacy-preserving continual learning and split learning respectively, demonstrating the effectiveness of these approaches in

Fourth International Conference on Computer Graphics, Image, and Virtualization (ICCGIV 2024), edited by Hoshang Kolivand, Ata Jahangir Moshayedi, Proc. of SPIE Vol. 13288, 132880W · © 2024 SPIE · 0277-786X · doi: 10.1117/12.3044867

handling sensitive medical data while maintaining model adaptability and collaborative training efficiency. Additionally, the work [4, 5] have emphasized the utility of AutoML tools and multimodal metadata augmentation in enhancing classification accuracy. In the operational domain, [6] highlight the implementation of resilience-aware MLOps, optimizing AI-based diagnostic systems. Also, Shen introduced a method to attack image segmentation[7]. Moreover, [8, 9] have contributed novel perspectives on deep learning methodologies, with Lei et al. introducing BP-CapsNet for diverse medical diagnostic applications and Zhang et al. advancing the field of AutoML by unifying data augmentation and neural architecture search.

Recent advancements in deep learning have significantly influenced medical image classification, notably improving automated disease diagnosis and enhancing medical research. Cai provided a thorough review of deep learning applications in medical image classification, emphasizing their role in enhancing diagnostic accuracy across various imaging modalities [10]. Yang introduced "MedMNIST v2," a benchmark for biomedical image classification that underscores the importance of high-quality datasets in training effective models [11]. Furthermore, [12, 13] have both highlighted the critical architectural choices in neural networks and their impacts on segmentation and classification outcomes. Kim demonstrated how transfer learning could be leveraged in medical image classification to adapt pre-trained models to new, smaller datasets, a technique beneficial in scenarios of data scarcity [14]. Lastly, Chen discussed the clinical implications of deep learning in medical image analysis, noting the necessity for continual algorithmic enhancement to manage diverse imaging tasks effectively [15]. Together, these studies underscore the rapid evolution and application of deep learning techniques in the medical field, each building upon the previous works to advance medical diagnostics.

3. METHOD

3.1 Curvature regularization

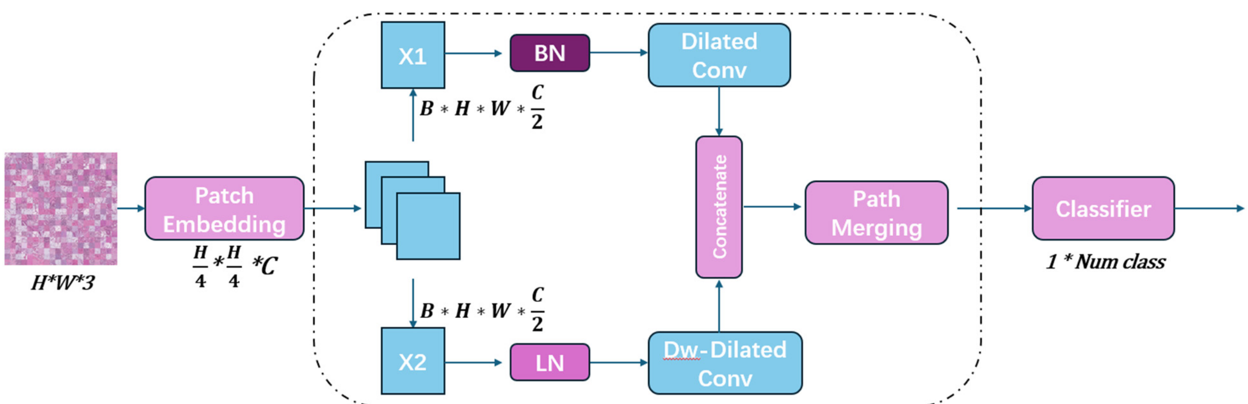

Figure 1. Overall architectural of MedMamba with Dilated Convolution: the framework encapsulated within the bounding box may be iterated several times.

To enhance the robustness of our medical image classification model, we employ Curvature Regularization, a technique designed to stabilize the learning process by controlling the magnitude of the Hessian's eigenvalues [17]. This regularization approach helps in making the model less sensitive to variations in input data, which is particularly beneficial in medical imaging where precision is crucial. The regularization term, L_r, is defined by $L_r = \sum_i p(\lambda_i)$ where λ_i denotes the eigenvalues of H and $p(t) = t^2$. This choice equates L_r to the Frobenius norm of H :

$$L_r = \text{trace}\big(p(H)\big) = E[z^T p(H) z] = E \parallel Hz \parallel^2 \tag{1}$$

where $z \sim \mathcal{N}(0, I_d)$, indicating a standard multivariate normal distribution. This formulation captures the expected squared norm of Hz, emphasizing the model's sensitivity to variations in different directions within the input space. The practical implementation simplifies L_r to:

$$L_r = \frac{1}{h^2} E \parallel \nabla \ell(x + hz) - \nabla \ell(x) \parallel^2 \tag{2}$$

where h is the discretization step size, adjusting the scale at which gradient variations are penalized. This expression aids in smoothening the gradient field locally. We select directions of known high curvature based on empirical analysis. We set z as $\operatorname{sign}(\nabla\ell(x))/\| \operatorname{sign}(\nabla\ell(x)) \|$, which prioritizes directions where the gradient is significant, further refining the focus of curvature regularization. This is incorporated into the loss function $\ell' = \ell + \gamma L_r$, where γ is a tuning parameter that balances the original loss and the regularization term, effectively enhancing the model's robustness against input perturbations.

3.2 Architecture of MedMamba with Dilated Convolution

We show our overall architectural in the Figure 1. We do curvature regularization to all the data, then we use Medmamba with Dilated Convolution to process all data. MedMamba's architecture [16] integrates dilated convolutions to enhance detail capture without increasing computational complexity. This adjustment targets the convolutional layers within the SS-Conv-SSM Block, broadening their receptive fields to better process spatial hierarchies in medical images.

Patch Embedding and Normalization: The input image $x \in \mathcal{R}^{H\times W\times 3}$ is divided into 4×4 non-overlapping patches and transformed into embeddings $x' \in \mathcal{R}^{\frac{H}{4}\times\frac{W}{4}\times C}$, with normalization applied subsequently.

Incorporation of Dilated Convolutions: The convolutional layers are adapted to use dilated convolutions, defined as:

$$y_{\text{conv}} = \text{ReLU}(Conv_d(x')) \tag{3}$$

where Conv_d denotes a convolution with a dilation rate d, optimizing spatial coverage and detail preservation.

Dual-branch SS-Convd-SSM Block: The SS-Convd-SSM Block, pivotal to MedMamba, operates in dual branches:

- Convd-branch: Applies dilated convolutions for enhanced local feature extraction, significantly benefiting from the dilated strategy.
- SSM-branch: Processes inputs through layer normalization and a linear layer, followed by the 2D Selective Scan module and depthwise separable convolutions. Feature Fusion and Output: The outputs of both branches are combined and undergo a channel shuffle to enhance feature integration:

$$y = \text{ChannelShuffle}(\text{Merge}\,(y_{\text{convd}}, y_{\text{SSM}})) \tag{4}$$

This architectural refinement enables MedMamba to capture intricate details in medical images, improving diagnostic accuracy through a more effective feature representation.

4. EXPERIMENTS

4.1 Dataset

We evaluate our proposed model using the PATHMNIST dataset, a subset of the MEDMNIST V2 [18], specifically designed for pathological image classification. It comprises pre-processed images of various pathologies, standardized to 28 * 28 pixels to facilitate uniform training and testing conditions. The dataset includes diverse classification tasks, ranging from binary to multiclass scenarios, with a significant volume of samples adequate for both training and rigorous evaluation.

4.2 Implementation Details

The architecture adaptation for PATHMNIST includes dilated convolutions to enhance detail recognition in pathological images. We optimize the network depth and hyperparameters based on preliminary evaluations. Similar to procedures used in MEDMNIST evaluations [18], we employ a layer-depth configuration optimized for 2D images, ensuring efficient processing and minimal information loss. Following the official MEDMNIST guidelines, the model is trained and tested using standard cross-entropy loss for classification tasks. The performance is assessed using the established metrics of accuracy (ACC) and the area under the receiver operating characteristic curve (AUC). These metrics are crucial for evaluating the effectiveness of medical image classification models, calculated as:

$$\text{ACC} = \frac{TP + TN}{TP + FP + TN + FN} \tag{5}$$

where TP, TN, FP, and FN denote the true positive, true negative, false positive, and false negative rates, respectively.

$$\mathrm{AUC} = \int_0^1 \mathrm{ROC}(r)\, dr \tag{6}$$

The ROC (r) represents the receiver operating characteristic curve as a function of the false positive rate, and dr denotes the differential element of the rate.

To ascertain the robustness of our model on the PATHMNIST dataset, we conducted adversarial testing using Fast Gradient Sign Method (FGSM) and Projected Gradient Descent (PGD). We compared our results with popular architectures like ResNet-18, ResNet-50, and automated machine learning methods including Auto-sklearn and Google AutoML. The adversarial tests evaluate the model's performance under potential real-world attack scenarios, ensuring the model's practical applicability in medical image analysis.

We implemented FGSM and PGD attacks to generate adversarial examples. For FGSM, we used:

$$x_{\mathrm{adv}} = x + \epsilon \cdot \mathrm{sign}\big(\nabla_x J(\theta, x, y)\big) \tag{7}$$

where ϵ is the perturbation magnitude. For PGD, we iteratively apply:

$$x^{(t+1)} = \mathrm{Proj}_{x+S}\left(x^{(t)} + \alpha \cdot sign\left(\nabla_x J\big(\theta, x^{(t)}, y\big)\right)\right) \tag{8}$$

where α is the step size, and Proj is the projection onto the set of allowable solutions S.

4.3 Result

The evaluation of our model focuses on its robustness against prevalent adversarial attacks, specifically the Fast Gradient Sign Method (FGSM) and Projected Gradient Descent (PGD). We compare its performance with well-established models, including traditional architectures like ResNet-18 and ResNet-50, as well as AutoML solutions such as Auto-sklearn and Google AutoML. This comprehensive comparison is critical for understanding the model's efficacy in adversarial scenarios, particularly in the context of medical image analysis using the PATHMNIST dataset.

Table 1. Comparison of model performance in terms of ACC under clean and attack scenarios.

Model	Clean ACC	FGSM ACC	PGD ACC	Avg Drop (%)
Our Model	0.90	0.85	0.84	5.6
ResNet-18	0.90	0.84	0.82	8.9
ResNet-50	0.91	0.85	0.83	8.8
Auto-sklearn	0.72	0.65	0.63	12.5
AutoML	0.73	0.66	0.64	12.3

In table 1, our model demonstrates superior robustness, with an average accuracy drop of only 5.6% under adversarial conditions, significantly less than the traditional ResNet architectures and AutoML solutions. The ResNet-18 and ResNet-50 models exhibit an average drop of 8.9% and 8.8% respectively, indicating their lesser resilience to adversarial perturbations. The AutoML models, Auto-sklearn and Google AutoML, suffer even higher drops of 12.5% and 12.3%, respectively. This substantial performance degradation is attributed to their general optimization strategies, which are not fine-tuned for the small-scale, detailed features inherent in the 28×28 pixel PATHMNIST images. These results underscore the necessity for specialized models tailored to specific dataset characteristics, particularly when robustness against adversarial attacks is critical.

Table 2. Comparison of model performance in terms of AUC under clean and attack scenarios.

Model	Clean AUC	FGSM AUC	PGD AUC	Avg Drop (%)
Our Model	0.98	0.94	0.93	4.1
ResNet-18	0.98	0.94	0.92	6.1
ResNet-50	0.99	0.95	0.93	6.1
Auto-sklearn	0.95	0.88	0.86	9.5
AutoML	0.94	0.87	0.85	9.6

In terms of AUC result in Table 2, our model again exhibits greater robustness with an average drop of only 4.1%. In comparison, the ResNet-18 and ResNet-50 models both show a drop of 6.1%. Auto-sklearn and Google AutoML experience even higher drops of 9.5% and 9.6%, respectively. These findings further validate that our model's architecture is better suited to handle adversarial attacks while maintaining high performance. This is particularly crucial in medical imaging applications, where even minor inaccuracies can lead to significant diagnostic errors.

In summary, the comparative analysis highlights that our model outperforms both traditional and AutoML architectures under adversarial conditions. Its tailored design for the PATHMNIST dataset's specific characteristics ensures minimal performance degradation, thereby enhancing its reliability for critical applications such as medical image analysis. The results clearly illustrate the importance of developing specialized architectures to effectively counter adversarial attacks and maintain high diagnostic accuracy.

5. DISCUSSION AND CONCLUSION

In this study, we have demonstrated the effectiveness and robustness of our proposed model using the PATHMNIST dataset, with a particular focus on adversarial resilience. Our model not only outperformed traditional architectures like ResNet-18 and ResNet-50 but also showed superior robustness compared to AutoML solutions such as Auto-sklearn and Google AutoML, especially under adversarial attack scenarios. The findings highlight the importance of designing specialized neural network architectures that are optimized for the specific challenges posed by medical image datasets, which often include lower resolution and highly variable image qualities.

In conclusion, our model's ability to maintain high accuracy and robustness under adversarial conditions suggests that it is well-suited for deployment in critical medical imaging tasks. By sustaining less performance degradation under attacks like FGSM and PGD, our model demonstrates a significant potential for reliable use in clinical environments where diagnostic precision is paramount. This study reinforces the critical role of specialized architectures in enhancing the robustness and accuracy of medical image analysis. Future work will involve exploring further optimizations and extensions of our approach to ensure even greater resilience and performance across a wider array of medical imaging datasets and adversarial scenarios.

REFERENCES

[1] Radford, Alec, et al. "Learning transferable visual models from natural language supervision." International conference on machine learning. PMLR, 2021.

[2] Verma, Tanvi, et al. "Privacy-preserving continual learning methods for medical image classification: a comparative analysis." Frontiers in Medicine 10 (2023).

[3] Li, Zhuohang, et al. "Split Learning for Distributed Collaborative Training of Deep Learning Models in Health Informatics." AMIA Annual Symposium Proceedings. Vol. 2023. American Medical Informatics Association, 2023.

[4] Yang, Jiancheng, et al. "Medmnist v2-a large-scale lightweight benchmark for 2d and 3d biomedical image classification." Scientific Data 10.1 (2023): 41.

[5] Gordienko, Yuri, Maksym Shulha, and Sergii Stirenko. "Effect of Multimodal Metadata Augmentation on Classification Performance in Deep Learning." International Conference on Trends in Sustainable Computing and Machine Intelligence. Singapore: Springer Nature Singapore, 2023.

[6] Moskalenko, Viacheslav, and Vyacheslav Kharchenko. "Resilience-aware MLOps for AI-based medical diagnostic system." Frontiers in Public Health 12 (2024): 1342937.

[7] Shen, Yifan, Zhengyuan Li, and Gang Wang. "Practical Region-level Attack against Segment Anything Models." arXiv preprint arXiv:2404.08255 (2024).

[8] Lei, Yongjia, et al. "BP-CapsNet: An image-based Deep Learning method for medical diagnosis." Applied Soft Computing 146 (2023): 110683.

[9] Zhang, Jianwei, et al. "Auto machine learning for medical image analysis by unifying the search on data augmentation and neural architecture." arXiv preprint arXiv:2207.10351 (2022).

[10] Cai, Lei, Jingyang Gao, and Di Zhao. "A review of the application of deep learning in medical image classification and segmentation." Annals of translational medicine 8.11 (2020).

[11] Yang, Jiancheng, et al. "Medmnist v2-a large-scale lightweight benchmark for 2d and 3d biomedical image classification." Scientific Data 10.1 (2023): 41.
[12] Malhotra, Priyanka, et al. "Deep neural networks for medical image segmentation." Journal of Healthcare Engineering 2022 (2022).
[13] Suganyadevi, S., V. Seethalakshmi, and K. Balasamy. "A review on deep learning in medical image analysis." International Journal of Multimedia Information Retrieval 11.1 (2022): 19-38.
[14] Kim, Hee E., et al. "Transfer learning for medical image classification: a literature review." BMC medical imaging 22.1 (2022): 69.
[15] Chen, Xuxin, et al. "Recent advances and clinical applications of deep learning in medical image analysis." Medical Image Analysis 79 (2022): 102444.
[16] Moosavi-Dezfooli, Seyed-Mohsen, et al. "Robustness via curvature regularization, and vice versa." Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition. 2019.
[17] Yue, Yubiao, and Zhenzhang Li. "Medmamba: Vision mamba for medical image classification." arXiv preprint arXiv:2403.03849 (2024).
[18] Yang, Jiancheng, et al. "Medmnist v2-a large-scale lightweight benchmark for 2d and 3d biomedical image classification." Scientific Data 10.1 (2023): 41.